



Prediction Intervals for Regression: *Theory, Methods, and Applications*

Master Thesis Defense

Marie Køhler Gudum
23 June 2026



Motivation

Imagine you are selling your house in Frederiksberg.

Motivation

Imagine you are selling your house in Frederiksberg.
A model gives you a number:

$$\hat{y} = 5.0 \text{ mio. DKK}$$

Motivation

Imagine you are selling your house in Frederiksberg.
A model gives you a number:

$$\hat{y} = 5.0 \text{ mio. DKK}$$

But is the true value likely to be between 4.9-5.1, or could it be anywhere in 3.5-6.5?

Motivation

Imagine you are selling your house in Frederiksberg.
A model gives you a number:

$$\hat{y} = 5.0 \text{ mio. DKK}$$

But is the true value likely to be between 4.9-5.1, or could it be anywhere in 3.5-6.5?

This is important in decision-making with high risk at stake:

- Mortgage lending decisions
- Insurance risk assessment
- Automated valuation models

Motivation

Imagine you are selling your house in Frederiksberg.
A model gives you a number:

$$\hat{y} = 5.0 \text{ mio. DKK}$$

But is the true value likely to be between 4.9-5.1, or could it be anywhere in 3.5-6.5?

This is important in decision-making with high risk at stake:

- Mortgage lending decisions
- Insurance risk assessment
- Automated valuation models

A point estimate is not enough. Decision-making requires uncertainty quantification.

Outline

1. Notation and definitions
2. Different methods
3. Empirical results
4. Main conclusions

What is a prediction interval

Construct an interval $C(X_{n+1})$ such that

$$\mathbb{P}(Y_{n+1} \in C(X_{n+1})) \geq 1 - \alpha$$

Example:

Point estimate

$$\hat{y} = 5.0$$

Prediction interval



$$[4.4, 5.6]$$

What do we want?

Validity

The interval should contain the future observation with probability at least $1 - \alpha$.

Efficiency

The interval should be informative and therefore as short as possible.

Adaptivity

The interval width should vary with the level of uncertainty.

Can we achieve all three simultaneously?

What does it mean to be valid?

Marginal Validity

$$\mathbb{P}(Y \in C(X)) \geq 1 - \alpha$$

Guarantee on average.

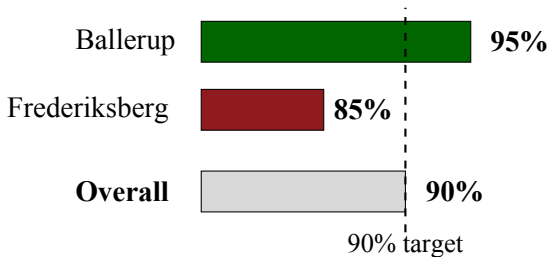
Conditional Validity

$$\mathbb{P}(Y \in C(X) \mid X = x) \geq 1 - \alpha$$

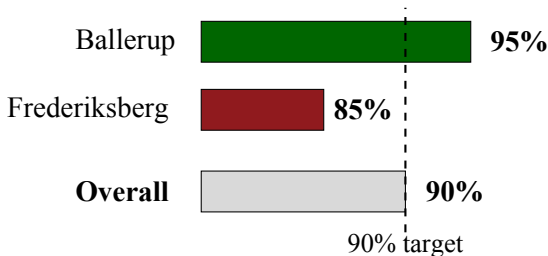
Guarantee for every x .

These are fundamentally different.

Why marginal validity may be misleading

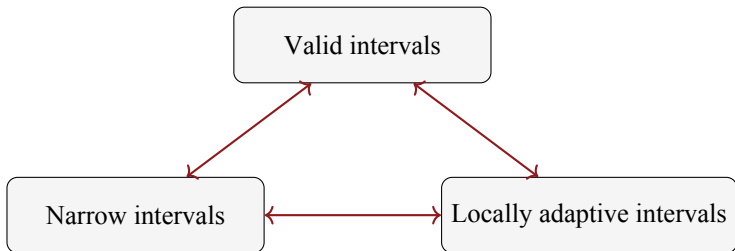


Why marginal validity may be misleading

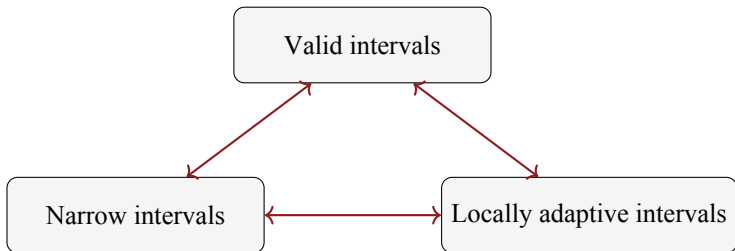


Good average performance
does not imply good local performance.

The fundamental challenge



The fundamental challenge



Cannot generally be achieved simultaneously

Three approaches to the problem

Three approaches to the problem

Gaussian Prediction Intervals

Strong assumptions in exchange for strong theoretical guarantees.

Three approaches to the problem

Gaussian Prediction Intervals

Strong assumptions in exchange for strong theoretical guarantees.

Quantile Regression

Flexible estimation of conditional quantiles and adaptive interval widths.

Three approaches to the problem

Gaussian Prediction Intervals

Strong assumptions in exchange for strong theoretical guarantees.

Quantile Regression

Flexible estimation of conditional quantiles and adaptive interval widths.

Conformal Prediction

Distribution-free validity under minimal assumptions.

Classical Gaussian prediction intervals

Assume $Y = X^\top \beta + \varepsilon$ with $\varepsilon \stackrel{\text{i.i.d.}}{\sim} N(0, \sigma^2)$ and $\varepsilon \perp\!\!\!\perp X$

Classical Gaussian prediction intervals

Assume $Y = X^\top \beta + \varepsilon$ with $\varepsilon \stackrel{\text{i.i.d.}}{\sim} N(0, \sigma^2)$ and $\varepsilon \perp\!\!\!\perp X$

Strengths

- Exact finite-sample validity
- Conditional validity
- Statistically efficient

Classical Gaussian prediction intervals

Assume $Y = X^\top \beta + \varepsilon$ with $\varepsilon \stackrel{\text{i.i.d.}}{\sim} N(0, \sigma^2)$ and $\varepsilon \perp\!\!\!\perp X$

Strengths

- Exact finite-sample validity
- Conditional validity
- Statistically efficient

Weaknesses

- Linearity
- Homoskedasticity
- Gaussian errors

Quantile Regression

Idea: estimate $q_{\alpha/2}(x)$ and $q_{1-\alpha/2}(x)$ directly.

Quantile Regression

Idea: estimate $q_{\alpha/2}(x)$ and $q_{1-\alpha/2}(x)$ directly.

Strengths

- Adapts to heteroskedasticity
- Allows asymmetric intervals
- Flexible extension via QRF

Quantile Regression

Idea: estimate $q_{\alpha/2}(x)$ and $q_{1-\alpha/2}(x)$ directly.

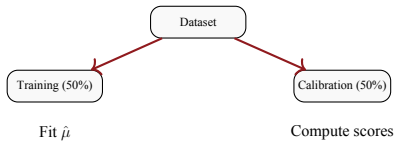
Strengths

- Adapts to heteroskedasticity
- Allows asymmetric intervals
- Flexible extension via QRF

Weaknesses

- Performance depends on model quality
- No finite-sample guarantee

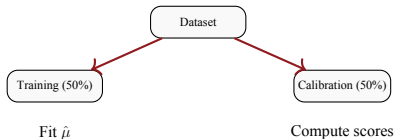
(Split) Conformal Prediction



(Split) Conformal Prediction

Key guarantee:

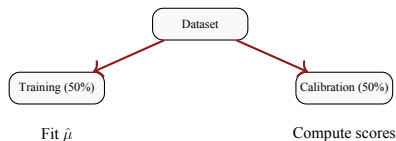
$$\mathbb{P}(Y_{n+1} \in C(X_{n+1})) \geq 1 - \alpha$$



(Split) Conformal Prediction

Key guarantee:

$$\mathbb{P}(Y_{n+1} \in C(X_{n+1})) \geq 1 - \alpha$$



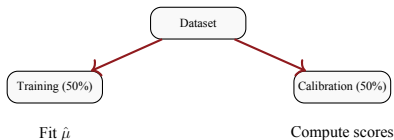
Strengths

- Finite-sample validity
- Distribution-free
- Independent of the prediction model

(Split) Conformal Prediction

Key guarantee:

$$\mathbb{P}(Y_{n+1} \in C(X_{n+1})) \geq 1 - \alpha$$



Strengths

- Finite-sample validity
- Distribution-free
- Independent of the prediction model

Weaknesses

- Marginal validity only
- Limited local adaptivity
- Efficiency depends on model quality

How stable is conformal coverage?

For split conformal prediction, the training-conditional coverage

$$\mathbb{P}(Y_{n+1} \in C(X_{n+1}) \mid \mathcal{D}_n)$$

is random.

Under suitable assumptions,

$$\mathbb{P}(Y_{n+1} \in C(X_{n+1}) \mid \mathcal{D}_n) \sim \text{Beta}((1 - \alpha)(n + 1), \alpha(n + 1))$$

Consequences:

- Mean = $1 - \alpha$
- Variance = $\frac{\alpha(1 - \alpha)}{n + 2}$
- Coverage becomes more stable as n increases

The tradeoff

Method	Marg. Validity	Cond. Validity	Dist.-free	Adaptive Width	Finite Sample
Gaussian PI	(✓)	(✓)	×	×	(✓)
QR	≈	≈	×	✓	×
CP	✓	×	✓	×	✓
Mondrian CP	✓	×	✓	≈	✓
CQR	✓	×	✓	✓	✓

✓ = guaranteed property (✓) = guaranteed property under model assumptions
 × = not guaranteed ≈ = holds asymptotically or only approximately

Simulation study

What happens when assumptions fail?

Simulation study

What happens when assumptions fail?

Scenarios:

- Linear vs nonlinear regression
- Symmetric vs asymmetric errors
- Homoskedastic vs heteroskedastic noise

Simulation study

What happens when assumptions fail?

Scenarios:

- Linear vs nonlinear regression
- Symmetric vs asymmetric errors
- Homoskedastic vs heteroskedastic noise

Methods compared:

Gaussian, QR, RQ, CP, and CQR.

Underlying prediction models:

Linear model and random forests

Four experiments, four challenges

Exp. 1: Baseline

$Y = \beta X + a + \varepsilon, \varepsilon \sim N(0, \sigma^2)$
All assumptions hold.

Exp. 2: Asymmetric noise

$\varepsilon \sim \text{Exp}(1/\sigma) - \sigma$
Tests symmetry assumptions.

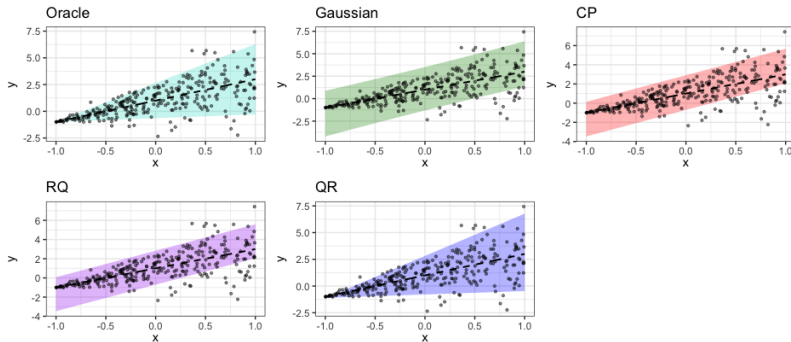
Exp. 3: Nonlinear mean

$\mu(x) = \sin(2\pi x), \varepsilon \sim N(0, \sigma^2)$
Tests model flexibility.

Exp. 4: Heteroskedastic

$\sigma(x) = 1 + x$
Tests local adaptivity.

Why global calibration fails locally



- Oracle width grows linearly with x - Gaussian, CP, and RQ stay flat
- Flat intervals overcover for small x , undercover for large x
- QR adapts in width

Heatmap over results - with CQR

Greener = closer to nominal coverage/oracle length.
No method is green everywhere.



Main results

Gaussian

performs well when
assumptions hold

Conformal

preserves coverage
across all settings

QR / QRF

improve adaptivity,
less stable coverage

**No method dominates across all settings.
Performance depends on the data-generating process**

The limitations

- Gaussian intervals fail under misspecification
- Quantile methods fail when quantiles are poorly estimated
- Conformal intervals fail to adapt to local variation

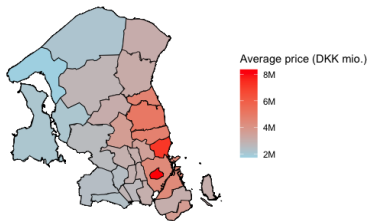
Case study: Housing prices

Data:

- 227,515 housing transactions
- Capital Region of Denmark, 2000-2025

Challenges:

- Nonlinearity
- Heteroskedasticity
- Local variation



Average price (DKK mio.) across municipalities

Case study: summary across methods

Average marginal coverage and interval width, $B = 1000$ repeated splits.

Method	Coverage	Width (mio. DKK)
Gaussian (LM)	0.911	3.7
Global CP (RF)	0.900	2.6
Mondrian CP (area) (RF)	0.900	2.6
QR (RF)	0.950	3.5
<i>Nominal target</i>	<i>0.900</i>	-

- All methods are valid on average
- QR overcovers, paying for its local stability with wider intervals
- CP remains narrowest, despite weaker local adaptivity

Main lessons

1. Validity is not a single concept
2. Conditional validity is fundamentally difficult
3. No method simultaneously optimizes validity, efficiency and adaptivity
4. Method choice depends on the objective

Limitations & future work

Limitations

- QRF performance is sensitive to tuning
- Same train/calibration split used across all methods, for comparability

Limitations & future work

Limitations

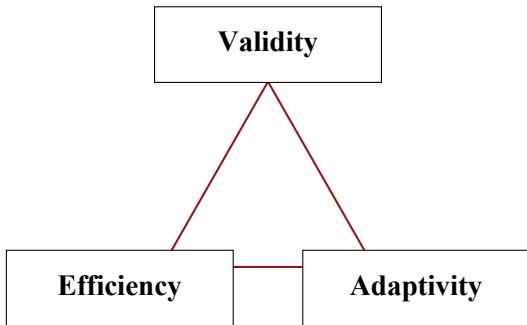
- QRF performance is sensitive to tuning
- Same train/calibration split used across all methods, for comparability

Future work

- Other base learners: boosting, neural nets
- Cross-validation methods
- Adaptive train/calibration split ratios

Conclusion

The construction of prediction intervals is fundamentally a trade-off problem.



Thank you for listening!